

Master in Artificial Intelligence

Advanced Human Language Technologies

Neural
Networks
NERC

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables



UNIVERSITAT POLITÈCNICA DE CATALUNYA
BARCELONATECH

Facultat d'Informàtica de Barcelona



Outline

1 Neural Networks NERC

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks
NERC

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Task 1.2 - NERC using neural networks

Assignment

Write a python program that parses the given XML and recognizes and classifies drug names. The program must use a neural network approach.

```
$ python3 ./nn-NER.py devel.xml result.out
```

```
DDI-DrugBank.d278.s0|0-9|Enoxaparin|drug
```

```
DDI-DrugBank.d278.s0|93-108|pharmacokinetics|group
```

```
DDI-DrugBank.d278.s0|113-124|eptifibatide|drug
```

```
DDI-MedLine.d88.s0|15-30|chlordiazepoxide|drug
```

```
DDI-MedLine.d88.s0|33-43|amphetamine|drug
```

```
DDI-MedLine.d88.s0|49-55|cocaine|drug
```

```
DDI-MedLine.d88.s1|82-95|benzodiazepine|drug
```

```
...
```

Outline

1 Neural Networks NERC

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks
NERC

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

General Structure

The general structure is basically the same than for the traditional ML approach:

- B-I-O schema
- Two programs: one learner and one classifier.
- The learner loads the training (Train) and validation (Devel) data, formats/encodes it appropriately, and feeds it to the model, together with the ground truth.
- The classifier loads the test data, formats/encodes it in the same way that was used in training, and feeds it to the model to get a prediction.

In the case of NN, we don't need to extract features (though we **do need** proper input encoding) normalmente tienes 5000 neuronas y un vector de input de 5000 de manera que solo 1 es 1 y el resto 0
we need to encode the vector, but we do not need feature

Input Encoding

- The input/output layers of a NN are vectors of neurons, each set to 0/1.
- Modern deep learning libraries handle this in the form of *indexes* (i.e. just provided the *position* of active neurons, omitting zeros). *si es 26 significa que es todo 0 menos en el indice 26*
- For instance, in a LSTM, each input word in the sequence may be encoded as the concatenation of different vectors each containing information about some aspect of the word (form, lemma, PoS, suffix...)
- Each vector will have only one active neuron (*one-hot-encoding*), indicated by its *index*. This input is usually fed to an embedding layer.
- Our learner will need to create and store *index* dictionaries to be able to map the index number assigned to each word, label, or any other used piece of information. See class *Codemaps*.

Outline

1 Neural Networks NERC

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks
NERC

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Outline

1 Neural Networks NERC

2 General Structure

3 Detailed Structure

- Learner
 - Classifier
 - Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks
NERC

General
Structure

Detailed
Structure
Learner

Core task

Goals &
Deliverables

Learner - Main program

```
1 def train(trainfile, validationfile, params, modelname) :
2     """
3     learns a NN model using trainfile as training data, and validationfile
4     as validation data. Saves learnt model in a file named modelname
5     """
6     # load train and validation data in a suitable form
7     traindata = Dataset(trainindir)
8     valdata = Dataset(validationindir)
9
10    # create indexes from training data
11    codes = Codemaps(traindata, params)
12    # encode datasets
13    train_loader = encode_dataset(traindata, codes, params)
14    val_loader = encode_dataset(valdata, codes, params)
15
16    # build network
17    network = nercLSTM(codes)
18
19    # save indexes
20    os.makedirs(modelname, exist_ok=True)
21    codes.save(modelname+"/codemaps")
22
23    # train network, keep the best performing model
24    best = 0
25    for epoch in range(params['epochs']):
26        acc = train(network, epoch, train_loader)
27        if acc>best :
28            best = acc
29        torch.save(network, os.path.join(modelname, f"network.nn"))
```

Neural
Networks
NERC

General
Structure

Detailed
Structure

Learner

Core task

Goals &
Deliverables

Outline

1 Neural Networks NERC

2 General Structure

3 Detailed Structure

- Learner
- **Classifier**
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks
NERC

General
Structure

Detailed
Structure
Classifier

Core task

Goals &
Deliverables

Classifier - Main program

```
1 def predict(modelname, datafile, params, outfile) :  
2     '''  
3     Loads a NN model from 'modelname' and uses it to extract drugs  
4     in datafile. Saves results to 'outfile' in the appropriate format.  
5     '''  
6     # Load model  
7     model = torch.load(os.path.join(modelname, "network.nn"),  
8                         map_location=torch.device(used_device))  
9     model.eval()  
10    # load indexes  
11    codes = Codemaps(os.path.join(modelname, "codemaps"), params)  
12    # load data to annotate  
13    testdata = Dataset(datafile)  
14    test_loader = encode_dataset(testdata, codes, params)  
15    # run each sentence through the NN, get results  
16    Y = []  
17    for X in test_loader:           you have for every sentence in the batch o dataset and for  
18        y = model.forward(*X)       each word in a sentence  
19        Y.extend([[codes.idx2label(torch.argmax(w)) for w in s] for s in y] )  
20                                     translates back the output  
21    # print results  
22    output_entities(testdata, Y, codes, outfile)
```

Neural
Networks
NERC

General
Structure

Detailed
Structure

Classifier

Core task

Goals &
Deliverables

Outline

1 Neural Networks NERC

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks
NERC

General
Structure

Detailed
Structure

Auxiliary classes

Core task

Goals &
Deliverables

Auxiliary classes - Dataset

```
1 class Dataset:
2     ## constructor: parses datafile XML file, tokenizes each sentence, and
3     ## stores a list of sentences, as well as ground truth for each token
4     def __init__(self, datafile)
5
6     ## iterator to get all sentences in the data set.
7     ## for each sentence returns a triplet (text, tokens, labels)
8     def sentences(self)
9
10    ## iterator to get ids for sentence in the data set
11    def sentence_ids(self)
12
13    ## get tokens for one given its id
14    def get_sentence_tokens(self, sid) :
15    ## get text for one sentence given its id
16    def get_sentence_text(self, sid) :
17    ## get labels for one sentence given its id
18    def get_sentence_labels(self, sid) :
19    , , ,
```

Neural
Networks
NERC

General
Structure

Detailed
Structure

Auxiliary classes

Core task

Goals &
Deliverables

Auxiliary classes - Codemaps

```
1 class Codemaps :
2     # Constructor: create code mapper either from training data, or
3     #               loading codemaps from given file.
4     #               If 'data' is a Dataset, and lengths are not None,
5     #               create maps from given data.
6     #               If data is a string (file name), load maps from file.
7     def __init__(self, data, params)
8     # Save created codemaps in file named 'name'
9     def save(self, name)
10    # Convert a Dataset into lists of word codes and suffix codes
11    # Adds padding and unknown word codes.
12    def encode_words(self, data)
13    # Convert the gold labels in given Dataset into a list of label codes.
14    # Adds padding
15    def encode_labels(self, data)
16    # get word index size
17    def get_n_words(self)
18    # get suf index size
19    def get_n_sufs(self)
20    # get label index size
21    def get_n_labels(self)
22    # get index for given word
23    def word2idx(self, w)
24    # get index for given suffix
25    def suff2idx(self, s)
26    # get index for given label
27    def label2idx(self, l)
28    # get label name for given index
29    def idx2label(self, i)
```

handles the translation

Neural
Networks
NERC

General
Structure

Detailed
Structure

Auxiliary classes

Core task

Goals &
Deliverables

Required functions - network.py

```
1 class nercLSTM(nn.Module):           podemos modificar los sizes dels embeddings
2     def __init__(self, codes):
3         super(nercLSTM, self).__init__()
4         # get sizes from index
5         n_words = codes.get_n_words()
6         n_sufs = codes.get_n_sufs()
7         n_labels = codes.get_n_labels()
8         # create embedding layers
9         embW_sz, embS_sz = 100, 50   we have the embeddings for the words and the suffixes
10        self.embW = nn.Embedding(n_words, embW_sz)
11        self.embS = nn.Embedding(n_sufs, embS_sz)
12        self.dropW = nn.Dropout(0.1)
13        self.dropS = nn.Dropout(0.1)
14        # create LSTM layer + final linear classification layer
15        lstm_in_sz, lstm_out_sz = embW_sz+embS_sz, 200
16        self.lstm = nn.LSTM(lstm_in_sz, lstm_out_sz,
17                             bidirectional=True, batch_first=True)
18        self.out = nn.Linear(2*lstm_out_sz, n_labels)
19        since the lstm is bidirectional the output is twice the size
20    def forward(self, w, s):           concatenates and links the layers
21        x = self.embW(w)               # apply embedding layers to input
22        y = self.embS(s)
23        x = self.dropW(x)              # apply dropout to embeddings output
24        y = self.dropS(y)
25        # concatenate embeddigns for word + suffix
26        x = torch.cat((x, y), dim=2)
27        x = self.lstm(x)[0]            # feed concatenated vector to LSTM
28        x = self.out(x)               # feed LSTM output to classification layer
29        return x
30
```

Neural
Networks
NERC

General
Structure

Detailed
Structure

Auxiliary classes

Core task

Goals &
Deliverables

Network architecture

Neural
Networks
NERC


General
Structure

Detailed
Structure

Auxiliary classes

Core task

Goals &
Deliverables



we could add more vectors
f.i. for the pos, lemmas...
añadir un vector de features sin embedding (ver siguiente)

the idea is that we can extend the network or try with different

Outline

1 Neural Networks NERC

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks
NERC

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Build a good NN-based drug NERC

Strategy: Experiment with different architectures and possibilities.
Some elements you can play with:

- Embedding layers dimensions
- Initializing word embeddings with available pretrained models
- Max length and suffix length values
- Number and/or size of LSTM layers
- Used optimizer, learning rate, batch size, ...
- Number and kind of layers or activation functions
- Additional input layers (maybe with embeddings). **Attention:**
This will require extending class Codemaps to handle the codes of added input layers.
 - different length suffixes and/or prefixes
 - PoS tags
 - lemmas
 - feature layer (with information about capitalization, dashes, presence in external resources, etc)

Neural
Networks
NERC

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

una layer que tenga
nuestras features, la
network no sabe
como son las
palabras si estan en
mayusculas o con
slash. No lo tenemos
que embed xq ya
son vectores de 1 y 0

Build a good NN-based drug NERC

Warnings:

- Neural Network training uses randomization, so different runs of the same program will produce different results. For repeatable results, use a random seed (and/or run the training several times).
- During training, *accuracy* on training and validation sets is reported. Those values are usually over 98%. However, this is due to the fact that most of the words have label “0” (non-drug). Accuracy values around 98% roughly correspond to F_1 values under 25%. To get a reasonable F_1 , validation set accuracy should reach about 99.5%.

To precisely evaluate how your model is doing, **do not rely** on reported accuracy: run the classifier on the development set and use the evaluator.

Outline

1 Neural Networks NERC

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks
NERC

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Exercise Goals

What you should do:

- Experiment with architecture variations.
- Experiment with different learning hyperparameters.
- Experiment with different input information
- Keep track of tried variants and parameter combinations.

What you should **NOT** do:

- Get insights about errors from `devel` dataset. You have `train` for that. `devel` is only used to evaluate the performance of a given configuration.
- Select architectures or hyperparameters based on system performance on `test` dataset. You have `devel` for that. `test` is only used to evaluate model generalization ability once the best configuration has been chosen.

Deliverables

PYTHON3 RUN.PY PARSE
PYTHON3 RUN.PY TRAIN
EN RUN.PY HAY INSTRUCCIONES

At the end of the NERC task, you will need to deliver a single report on the work carried out on NERC-ML, NERC-NN, and NERC-LLM systems

So, during the development and experimentation on NERC-NN:

- keep track of tried/discarded architectures/hyperparameters
- keep track of tried/discarded input information.
- Record obtained results in the different experiments, and compile the information you'll later need to elaborate the report.

one of the goals is comparing nn with classical ml, was it faster?? ...

Neural
Networks
NERC

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables